

In [37]:

```
# Import Pandas
import pandas as pd
from sklearn.feature_extraction.text import CountVectorizer

# Import linear_kernel
from sklearn.metrics.pairwise import linear_kernel
from sklearn.metrics.pairwise import cosine_similarity

corpus=[ 'Julie loves me more than Linda loves me',
         'Jane likes me more than Julie loves me',
         'harry likes kiwi fruit']

# Creating coun vectorizer object. Note: we have removed the stop words
vectorizer = CountVectorizer(stop_words='english')
vectors = vectorizer.fit_transform(corpus)
print('\n ---Corpus converted to term-frequency vector:--\n', vectors.toarray())

#Array mapping from feature integer indices to feature name or we can say the unique terms present in the corpus
print('\n ---Unique terms in the corpus:--- \n', vectorizer.get_feature_names())

#Compute similarity score
cosine_sim = cosine_similarity(vectors, vectors)
cosine_sim2 = linear_kernel(vectors, vectors)
print('\n--Cosine similarity using cosine similarity function:--\n',cosine_sim)
print('\n--Cosine similarity using Linear Kernel function:--\n',cosine_sim2)

# You can write your own function as well for computing cosine similarity like following:
"""
doc1=vectors.toarray()[0, :]
doc2=vectors.toarray()[1, :]
doc3=vectors.toarray()[2, :]
cos_sim_doc1_doc2 = dot(doc1, doc2)/(norm(doc1)*norm(doc2))

cos_sim_doc1_doc3 = dot(doc1, doc3)/(norm(doc1)*norm(doc3))

cos_sim_doc2_doc3 = dot(doc2, doc3)/(norm(doc2)*norm(doc3))

print(cos_sim_doc1_doc2)
print(cos_sim_doc1_doc3)
print(cos_sim_doc2_doc3)
"""
```

```

---Corpus converted to term-frequency vector:--
[[0 0 0 1 0 0 1 2]
 [0 0 1 1 0 1 0 1]
 [1 1 0 0 1 1 0 0]]

---Unique terms in the corpus:---
['fruit', 'harry', 'jane', 'julie', 'kiwi', 'likes', 'linda', 'loves']

--Cosine similarity using cosine similarity function:--
[[1.          0.61237244  0.          ]
 [0.61237244  1.          0.25         ]
 [0.          0.25         1.          ]]

--Cosine similarity using Linear Kernel function:--
[[6.  3.  0.]
 [3.  4.  1.]
 [0.  1.  4.]]

```

Out[37]:

```

'\ndoc1=vectors.toarray()[0, :]\ndoc2=vectors.toarray()[1, :]\ndoc3=vector
s.toarray()[2, :]\ncos_sim_doc1_doc2 = dot(doc1, doc2)/(norm(doc1)*norm(doc
2))\n\ncos_sim_doc1_doc3 = dot(doc1, doc3)/(norm(doc1)*norm(doc3))\n\ncos_
sim_doc2_doc3 = dot(doc2, doc3)/(norm(doc2)*norm(doc3))\n\nprint(cos_sim_d
oc1_doc2)\nprint(cos_sim_doc1_doc3)\nprint(cos_sim_doc2_doc3)\n'

```

You can see that document 1 and 2 are more similar. You will also notice that there is difference between the similarity score given by the `cosine_similarity()` and `linear_kernel()`. This happens because by default 'CountVectorizer' fit.transform function does not length normalize the data.

Q2. Repeat the above exercise but this time use tf-idf weights to convert the corpus?

In [2]:

```

# Import Pandas
import pandas as pd
#Import TfidfVectorizer from scikit-learn
from sklearn.feature_extraction.text import TfidfVectorizer

# Import linear_kernel
from sklearn.metrics.pairwise import linear_kernel
from sklearn.metrics.pairwise import cosine_similarity

corpus=[ 'Julie loves me more than Linda loves me',
         'Jane likes me more than Julie loves me',
         'harry likes kiwi fruit']

#Define a TF-IDF Vectorizer Object. Remove all english stop words such as 'the', 'a'
tfidf = TfidfVectorizer(stop_words='english')
#tfidf = TfidfVectorizer()

#Construct the required TF-IDF matrix by fitting and transforming the data
tfidf_matrix = tfidf.fit_transform(corpus)

#Output the shape of tfidf_matrix

print(tfidf_matrix.shape)
print(tfidf_matrix.toarray())
#Array mapping from feature integer indices to feature name.
#print(tfidf.get_feature_names())
# Compute the cosine similarity matrix
cosine_sim = linear_kernel(tfidf_matrix, tfidf_matrix)
cosine_sim2 = cosine_similarity(tfidf_matrix, tfidf_matrix)

print('\n--Cosine similarity using cosine similarity function:--\n',cosine_sim)
print('\n--Cosine similarity using Linear Kernel function:--\n',cosine_sim2)

```

```

(3, 8)
[[0.          0.          0.          0.38550292 0.          0.
  0.50689001 0.77100584]
 [0.          0.          0.60465213 0.45985353 0.          0.45985353
  0.          0.45985353]
 [0.52863461 0.52863461 0.          0.          0.52863461 0.40204024
  0.          0.          ]]

```

```
--Cosine similarity using cosine similarity function:--
```

```

[[1.          0.53182464 0.          ]
 [0.53182464 1.          0.18487962]
 [0.          0.18487962 1.          ]]

```

```
--Cosine similarity using Linear Kernel function:--
```

```

[[1.          0.53182464 0.          ]
 [0.53182464 1.          0.18487962]
 [0.          0.18487962 1.          ]]

```

You can see that document 1 and 2 are more similar. You will also notice that there is difference between the similarity score given by the `cosine_similarity()` and `linear_kernel()`. This happens because by default `TfidfVectorizer()` `fit.transform` function length normalizes the data.

Q3. Now build a complete recommender system for a real world dataset. Read the data in the file 'movies_metadata.csv'. Use 'overview' attribute to compute similarity between the movies. Finally list top 10 similar movies to 'Father of the Bride Part II'. Note: You can use some basic pre-processing techniques for example, removing rows with missing values etc.

In [39]:

```
# Import Pandas
import pandas as pd

#Import TfidfVectorizer from scikit-learn
from sklearn.feature_extraction.text import TfidfVectorizer

# Import linear_kernel
from sklearn.metrics.pairwise import linear_kernel
```

In [40]:

```
# Load Movies Metadata
metadata = pd.read_csv('dataset\movies_metadata.csv', low_memory=False)

# Print the first three rows
print(metadata.head(3))

#Print plot overviews of the first 5 movies.
print(metadata['overview'].head())
```

```

adult      belongs_to_collection      budget \
0 False    {'id': 10194, 'name': 'Toy Story Collection', ... 30000000
1 False                                          NaN  65000000
2 False    {'id': 119050, 'name': 'Grumpy Old Men Collect... 0

                                genres \
0 [{'id': 16, 'name': 'Animation'}, {'id': 35, '...
1 [{'id': 12, 'name': 'Adventure'}, {'id': 14, '...
2 [{'id': 10749, 'name': 'Romance'}, {'id': 35, ...

                                homepage      id      imdb_id original_language
e \
0 http://toystory.disney.com/toy-story      862      tt0114709      e
n
1                                          NaN      8844      tt0113497      e
n
2                                          NaN      15602      tt0113228      e
n

                                original_title      overview
... \
0      Toy Story      Led by Woody, Andy's toys live happily in his ...
...
1      Jumanji      When siblings Judy and Peter discover an encha...
...
2      Grumpier Old Men      A family wedding reignites the ancient feud be...
...

                                release_date      revenue runtime \
0      1995-10-30      373554033.0      81.0
1      1995-12-15      262797249.0      104.0
2      1995-12-22              0.0      101.0

                                spoken_languages      status \
0      [{'iso_639_1': 'en', 'name': 'English'}]      Released
1      [{'iso_639_1': 'en', 'name': 'English'}, {'iso...      Released
2      [{'iso_639_1': 'en', 'name': 'English'}]      Released

                                tagline      title      vi
deo \
0                                          NaN      Toy Story      Fa
lse
1      Roll the dice and unleash the excitement!      Jumanji      Fa
lse
2      Still Yelling. Still Fighting. Still Ready for...      Grumpier Old Men      Fa
lse

                                vote_average vote_count
0              7.7      5415.0
1              6.9      2413.0
2              6.5       92.0

[3 rows x 24 columns]
0      Led by Woody, Andy's toys live happily in his ...
1      When siblings Judy and Peter discover an encha...
2      A family wedding reignites the ancient feud be...
3      Cheated on, mistreated and stepped on, the wom...
4      Just when George Banks has recovered from his ...
Name: overview, dtype: object

```

In [41]:

```
#Replace NaN with an empty string
metadata['overview'] = metadata['overview'].fillna('')
```

In [42]:

```
#Define a TF-IDF Vectorizer Object. Remove all english stop words such as 'the', 'a'
tfidf = TfidfVectorizer(stop_words='english')
```

```
#Construct the required TF-IDF matrix by fitting and transforming the data
tfidf_matrix = tfidf.fit_transform(metadata['overview'])
```

```
#Output the shape of tfidf_matrix
print(tfidf_matrix.shape)
```

```
#Array mapping from feature integer indices to feature name.
print(tfidf.get_feature_names()[5000:5010])
```

```
(45466, 75827)
['avails', 'avaks', 'avalanche', 'avalanches', 'avallone', 'avalon', 'avan
t', 'avanthika', 'avanti', 'avaracious']
```

In [43]:

```
# Compute the cosine similarity matrix
cosine_sim = linear_kernel(tfidf_matrix, tfidf_matrix)

print('cosine similiarity matrix shape:', cosine_sim.shape)

#Construct a reverse map of indices and movie titles
indices = pd.Series(metadata.index, index=metadata['title']).drop_duplicates()

print(indices[:10])
```

```
cosine similiarity matrix shape: (45466, 45466)
```

```
title
```

Toy Story	0
Jumanji	1
Grumpier Old Men	2
Waiting to Exhale	3
Father of the Bride Part II	4
Heat	5
Sabrina	6
Tom and Huck	7
Sudden Death	8
GoldenEye	9

```
dtype: int64
```

In [44]:

```
# Function that takes in movie title as input and outputs most similar movies
def get_recommendations(title, cosine_sim=cosine_sim):
    # Get the index of the movie that matches the title
    idx = indices[title]

    # Get the pairwise similarity scores of all movies with that movie
    sim_scores = list(enumerate(cosine_sim[idx]))

    # Sort the movies based on the similarity scores
    sim_scores = sorted(sim_scores, key=lambda x: x[1], reverse=True)

    # Get the scores of the 10 most similar movies
    sim_scores = sim_scores[1:11]

    # Get the movie indices
    movie_indices = [i[0] for i in sim_scores]

    # Return the top 10 most similar movies
    return metadata['title'].iloc[movie_indices]
get_recommendations('Father of the Bride Part II')
```

Out[44]:

```
6793      Father of the Bride
6571                Kuffs
6306      North to Alaska
19801             Babbitt
34466    You're Killing Me
13611    The Magic of Méliès
5005                Wendigo
27974      I Start Counting
43887    George of the Jungle 2
7097      The Out of Towners
Name: title, dtype: object
```


Code credit

<https://www.datacamp.com/community/tutorials/recommender-systems-python>
(<https://www.datacamp.com/community/tutorials/recommender-systems-python>)

More on Text similarity

<https://medium.com/@adriensieg/text-similarities-da019229c894> (<https://medium.com/@adriensieg/text-similarities-da019229c894>)

More on recommender system

https://goodboychan.github.io/chans_jupyter/python/datacamp/natural_language_processing/2020/07/17/04-TF-IDF-and-similarity-scores.html#Cosine-similarity
(https://goodboychan.github.io/chans_jupyter/python/datacamp/natural_language_processing/2020/07/17/04-TF-IDF-and-similarity-scores.html#Cosine-similarity)



In []: